



Multi-Crop Plant Recognition Pipeline

End-to-end production-ready system: 139 plant species classifier
with Streamlit web interface and AI-powered agronomy guidance

Radhey Pareek

Problem Statement & Solution

Challenge

Farmers need rapid, accurate plant species identification with actionable agronomy guidance for diseases, remedies, climate requirements, and nutrients.

Solution

139-class deep learning classifier with robust inference pipeline featuring Test Time Augmentation (TTA), multi-crop averaging, and bilingual AI guidance via Gemini integration.

- ✓ **Key Innovation:** Production-optimized pipeline balancing accuracy with speed for modest hardware deployments




Technical Architecture Overview



Data Pipeline

Raw images \rightarrow RGB_224x224 preprocessing with ImageNet normalization and robust imread handling



Model Training

EfficientNet-B0 backbone with AdamW optimizer, cosine scheduling, and mixed precision training



Inference Engine

TTA + 5-crop averaging for robust predictions with configurable confidence thresholds



Web Application

Streamlit UI with camera/upload support and integrated Gemini AI agronomy guidance

Technology Stack & Dependencies

Deep Learning Core

- PyTorch + torchvision
- timm model zoo
- Mixed precision (AMP)

Data Processing

- Albumentations (fast augmentation)
- OpenCV + Pillow
- Hydra configuration

Production Interface

- Streamlit web UI
- Google Gemini API
- Minimal runtime deps

Rationale: [timm](#) provides battle-tested backbones, [Albumentations](#) offers fast transforms, and [Streamlit](#) accelerates demo-ready deployment.


Data Pipeline & Processing

Dataset Structure


```
data/raw/<class>/*.jpg  
↳ (prepare_dataset.py)  
data/RGB_224x224/  
=%% train/ (70%)  
=%% val/ (15%)  
5%% test/ (15%)
```

Key Features

- Robust `imread_rgb()` handles grayscale, RGB, RGBA
- Stable `class_to_idx` mapping preservation
- ImageNet mean/std normalization
- Training at 160-192px for speed optimization



- **Speed Optimization:** Smaller image sizes and reduced validation frequency (every 4 epochs) for faster training cycles



Model Configuration & Training Strategy

1

Architecture

Backbone: EfficientNet-B0 (default) / MobileNetV2
(speed profile)

Classes: 139 plant species

Input: 192×192 RGB images

2

Optimization

Optimizer: AdamW ($\text{lr}=1\text{e-}3$, $\text{wd}=0.01$)

Scheduler: Cosine annealing

Precision: Mixed precision + channels-last

3

Training Setup

Batch size: 48 | **Epochs:** 30

Augmentation: Flips, rotations, color jitter, dropout

Validation: Every 4 epochs for efficiency

Robust Inference Engine

Inference Model Features

- Auto-loads `weights/best.pt` with embedded config
- Preserves `class_to_idx` mapping
- Configurable TTA modes
- Multi-crop averaging support

TTA Strategies

- **Off:** Original image only
- **Fast:** + horizontal flip
- **Full:** + vertical flip

01

5-Crop Extraction

Center + 4 corners for robust framing coverage

02

TTA Application

Apply augmentations to each crop independently

03

Prediction Averaging

Ensemble predictions across all crops and augmentations

04

Confidence Thresholding

Auto-retry with stronger TTA if confidence < 0.35

Streamlit Web Application



Image Input

Supports file upload and real-time camera capture with preprocessing pipeline integration



Smart Inference

Cached model loading, multi-crop predictions, and adaptive TTA based on confidence scores



AI Guidance

Bilingual agronomy advice (English/Hindi) via Gemini API with structured output formatting

Green-themed responsive design with animated background and graceful fallbacks for offline scenarios

Production Results & Performance

Model Checkpoints

Best model: weights/best.pt

Validation: Every 4th epoch

Metrics: Embedded in
checkpoint

139

192

5

Plant Classes

Comprehensive multi-crop species coverage

Input Size

Optimized resolution for speed-accuracy balance

Crop Ensemble

Center + corners for robust predictions


30

Training Epochs

Efficient convergence with cosine scheduling

```
import torch
ckpt =
torch.load('weights/best.
pt')
print(ckpt['best_metric'])
```

- ✓ **Key Achievement:** Production-ready pipeline with minimal dependencies and embedded configuration for consistent deployment



Deployment & Future Roadmap

1 Current State

- 154-class classifier
- Streamlit web interface
- Gemini AI integration
- Production optimizations

2 Near Term

- Evaluation metrics script
- Early stopping & EMA
- ONNX export for edge
- Grad-CAM visualizations

3 Long Term

- 🎯 Multi-language expansion
- 🎯 Offline AI fallbacks
- 🎯 TensorRT optimization
- 🎯 Class imbalance handling

Ready to Deploy: streamlit run app.py with Gemini API key for immediate agricultural AI assistance