

# Multi-Crop Plant Recognition – Project Report

---

## 1) Executive Summary

- Goal: Fast and accurate plant species recognition (139 classes) with agronomy guidance.
  - Outcome: Streamlit web app + Gemini integration. Predicts species, provides tips, disease analysis, remedies, climate/season, regions, characteristics, and nutrients.
  - Speed: Epoch time reduced significantly via EfficientNet\_b0, img size 192, fewer validations, and warm-start training.
- 

## 2) Objectives

- Build end-to-end dataset → training → inference web app pipeline.
  - Optimize for fast training on CPU/GPU-limited hardware.
  - Add actionable guidance for farmers: treatments, home remedies, prevention, and climate information.
  - Keep code modular and documented for maintenance.
- 

## 3) Technology Stack

Area	Choice	Why
Deep Learning	PyTorch + timm	Reliable, rich model zoo, easy training/inference
Model	efficientnet_b0	Excellent speed/accuracy tradeoff for 139 classes
Data Augmentation	Albumentations	Fast, flexible transforms
Web UI	Streamlit	Rapid development and deployment
API/LLM	Google Gemini	Rich language output for agriculture insights
Config	YAML + OmegaConf	Clean, reproducible hyperparameters

---

## 4) Workflow Diagram

```
graph LR
    A[Raw Images data/raw/] --> B[Prepare dataset RGB-224x224]
    B --> C[Train 30 epochs EfficientNet_b0@192px]
    C --> D{Validate every 4 epochs}
    D -->|best improves| E[Save weights/best.pt]
    D --> F[Log Val Acc]
    E --> G[Streamlit Inference]
    G --> H[Gemini Analysis Species+Disease+Remedies+Climate+Regions+Nutrients]
```

---

## 5) Dataset & Preprocessing

- Source: `data/raw/` with 139 class folders (RGB images).
  - Split strategy: Per-class stratified split into `train/val/test` via `src/data/prepare_dataset.py`.
  - Normalization: ImageNet mean/std; image resized to 192x192 during training and inference.
- 

## 6) Model & Training Setup

Item	Setting	Rationale
Backbone	efficientnet_b0	Lightweight, fast, robust for 139 classes
Image size	192	Lower compute per batch
Epochs	30	Balanced speed and convergence
Batch size	48	Fewer iterations per epoch
Optimizer	AdamW	Stable with weight decay
LR schedule	Cosine, warmup=3	Smooth learning rate evolution
Regularization	Mixup/CutMix: OFF	Speed-focused profile
Validation cadence	val_every=2	Reduces idle time on validation
Warm-start	freeze_backbone_epochs=3	Train classifier head first

Key files: `configs/train.yaml`, `src/training/train.py`, `src/models/build.py`.

---

## 7) Issues Faced and Fixes

Problem	Symptom/Trace	Root Cause	Fix
Streamlit import error	ModuleNotFoundError: No module named 'src'	App executed from project root not on path	Add project root to <code>sys.path</code>
Inference top-k bug	TypeError: unhashable type: 'list'	topk tensors converted to nested lists	Flatten <code>topk</code> outputs before zipping
Class mismatch	Assertion error: ncls != cfg.num_classes	New dataset had 139 classes	Update <code>configs/train.yaml</code>
Slow epochs	~30 min/epoch initially	Heavy model/size and frequent validation	Switch to EfficientNet_b0, img=192
Resume mismatch	Old checkpoints with different class count	Attempted to resume with older head	Clean <code>weights/</code> before starting

## 8) Web App Features

- Upload or capture a plant image (camera).
  - Displays Top-5 species predictions and confidence.
  - Dynamic, species-aware.
  - Gemini panel: Uses Google Gemini to output disease/status, treatments, home remedies, prevention, climate/season, regions, characteristics, and nutrients.
  - Green-themed UI with animated background, wide layout, and two-column input area.
-

## 9) Gemini Integration

Item	Detail
Library API key	<code>google-generativeai</code> Loaded from function arg → env ( <code>GOOGLE_API_KEY/GEMINI_API_KEY</code> ) → <code>secrets/gemini_api_key.txt</code> → Streamlit secrets
Prompt	Structured prompt guiding sections: species, diseases, treatments, prevention, climate/season, regions, characteristics, nutrients
Code	<code>src/services/gemini.py</code> → <code>analyze_plant_with_gemini()</code>

Security note: avoid committing secrets. Prefer environment variables or `streamlit secrets` in production.

---

## 10) Results & Accuracy

- Checkpoints: `weights/best.pt` updated when validation improves; rolling `last_epochXXX.pt` for auto-resume.
  - Validation accuracy: printed to terminal every `val_every` epochs and on the last epoch. Use these logs to track accuracy trend.
  - Optional next step: add `eval.py` to compute Top-1/Top-5, per-class accuracy, and confusion matrix on test split.
- 

## 11) Reproducibility – Commands

Task	Command
Install deps	<code>pip install -r requirements.txt</code>
Prepare dataset	<code>python src/data/prepare_dataset.py --input data/raw --output data/RGB_224x224 --val_split 0.15 --test_split 0.15</code>
Train (full, fast)	<code>python -m src.training.train --config configs/train.yaml</code>
Quick experiment	<code>python -m src.training.train --config configs/train.yaml --epochs 10 --limit_train 100</code>
Run web app	<code>streamlit run src/app/streamlit_app.py</code>

---

## 12) Achievements

- End-to-end system with 139-class recognition and actionable agronomy insights.
  - Fast training mode with EfficientNet\_b0 and smart scheduling.
  - Stable web app with improved UI/UX and LLM assistance.
  - Robustness improvements: import path fix, inference fix, config validation.
- 

## 13) Future Work

- Add `eval.py` for detailed metrics and a dashboard in `reports/`.
  - Early stopping, EMA model, and optional RandAugment for higher accuracy.
  - Export to ONNX/TensorRT for lower-latency inference.
  - Add Grad-CAM visualization in the app for explainability.
  - Multi-lingual Gemini responses and offline tips fallback.
-